ML/DL Assignment
EagleView

Object Detection is one of the fundamental tasks in computer vision with applications ranging across medicine, robotics, and many others. In this task, you are given a dataset that consists of images containing people and cars. The goal of this task is to train a model that can localize and classify each instance of person and car as accurately as possible.

Dataset Details:

Number of images:	2239
Size:	750 MB (approx.)
Categories:	2 (person and car)
Annotation format:	COCO

Data directory structure:

```
data
  |_ annotations
       |_ bbox-annotations.json
  |_ images
       |_ image_000000001.jpg
       |_ image_000000002.jpg
       |_ ...
       ...
```

License: These images are taken from OpenImages. All annotations are licensed under CC BY 4.0 and the images have a CC BY 2.0 license*

Link to download the dataset: https://evp-ml-data.s3.us-east-2.amazonaws.com/ml-interview/openimages-personcar/trainval.tar.gz

Results Format:

We want you to develop and train a model on this dataset to detect person and cars in an image. You are free to use any open-source software and repository.

Deliverable:

1. A brief write-up about your solution explaining your assumption, approach, metrics, and other artifacts.
2. A link to your solution code base hosted in a GitHub repository.

*we make no representations or warranties regarding the license status of each image and you should verify the license for each image yourself.

Note:

1. The labels are already present in this dataset, thus there is no need to annotate the images from scratch.
2. Kindly try to keep your code readable and object-oriented, if possible, with proper method names.
3. Provide documentation in the GitHub README and don't upload heavy files/folders for instance – datasets and model weights.